

CSC263 – Problem Set 4

Remember to write your **full name** and **student number** prominently on your submission. To avoid suspicions of plagiarism: at the beginning of your submission, **clearly state any resources (people, print, electronic) outside of your group, the course notes, and the course staff, that you consulted.**

Remember that you are required to submit your problem sets as both LaTeX .tex source files and .pdf files. There is a 10% penalty on the assignment for failing to submit both the .tex and .pdf.

Due Mar. 27, 2025, 22:00; required files: ps4.pdf, ps4.tex, ps4.py

Answer each question completely, always justifying your claims and reasoning. Your solution will be graded not only on correctness, but also on clarity. Answers that are technically correct but are hard to understand will not receive full marks. Mark values for each question are contained in the [square brackets].

You may work in groups of up to TWO to complete these questions.

1. [12 points] You are given a weighted, connected, undirected graph $G = (V, E)$ and one of its minimum spanning trees $T \subseteq E$. Furthermore, you are guaranteed that the edge weights are all distinct.


- (a) [6 points] A new edge (u, v) with weight $w_{u,v}$ ($u, v \in V$) is added to G , resulting in a new graph $G' = (V, E \cup \{(u, v)\})$. How do you efficiently find a minimum spanning tree T' of G' ? Describe and justify your algorithm in concise and precise English, and analyse its runtime. To get full marks, your algorithm's worst-case runtime must be in $\mathcal{O}(|V|)$.

- (b) [4 points] Prove that the second lightest edge in the graph should belong to every minimum spanning tree.

Suppose for contradiction that there exists a minimum spanning tree that does not contain the second lightest edge. Assume that this edge connects some vertices u, v . It must be true that there is some path in the minimum spanning tree T that connects u and v . We make the assumption that G is simple (so there is not another edge directly connecting u, v), implying that this path traverses at least 2 edges. Of these edges, there is guaranteed to be one edge (r, s) that is strictly heavier than (u, v) . Create a new subgraph which contains all the edges of T , but contains the edge (u, v) and excludes the edge (r, s) . Notice that this subgraph is a spanning tree whose total edge weight is less than T , which contradicts the fact that T is a minimum spanning tree. Therefore all minimum spanning trees should contain the second lightest edge.

- (c) [2 points] Prove that the third lightest edge in the graph may not be in a minimum spanning tree.

Consider the graph



It is quite clear that the only minimum spanning tree for this graph is the subgraph obtained by removing the edge with weight 3. But this edge is also the third lightest edge, which illustrates what we wanted.

2. [6 points] You are given a $k \times k$ grid of squares (think of a chess board). Each square is identified by $(i, j) \in [k] \times [k]$ where $[k] = \{1, 2, \dots, k\}$, i is the row number, and j is the column number. From any square you are allowed to make one of the following moves: (i) m steps in the vertical direction and n steps in the horizontal direction, or (ii) n steps in the vertical direction and m steps in the horizontal direction. If you encounter the edge of the board (either a vertical edge or a horizontal edge) while making the move, then the move is invalid. For example, if $m = 1$ and $n = 2$, then this corresponds to the moves made by the knight on a chess board. You play this game with your friend. Your friend first chooses a square $(i, j) \in [k] \times [k]$. Your task is to find another square $(i', j') \in [k] \times [k]$ such that there is no valid sequence of moves from (i, j) to (i', j') , or declare that no such (i', j') exists. Devise an $O(k^2)$ algorithm to solve this task.

Solution: Our algorithm works as follows. First, we treat the grid as a graph, where each square is a vertex, and

$$((i, j), (i - n, j - m)), ((i, j), (i - n, j + m)), ((i, j), (i + n, j - m)), ((i, j), (i + n, j + m))$$

are edges, if the squares are in $[k] \times [k]$. Notice that in this graph, there are k^2 vertices and at most $4k^2$ edges. We perform a DFS on this graph originating from the original square chosen, which we call (i, j) . This has a run time of $\mathcal{O}(k^2)$. Afterwards, we iterate through each vertex with a nested for loop. If we find a vertex (i', j') that has not been visited, that implies that it cannot be reached from (i, j) , so we return that square. If we completely iterate through the grid without finding an unvisited square, then that means all squares are reachable, so no square is unreachable.

Programming Question

The best way to learn a data structure or an algorithm is to code it up. In each problem set, we will have a programming exercise for which you will be asked to write some code and submit it. You may also be asked to include a write-up about your code in the PDF/TeXfile that you submit. Make sure to **maintain your academic integrity** carefully, and protect your own work. The code you submit will be checked for plagiarism. It is much better to take the hit on a lower mark than risking much worse consequences by committing an academic offence.

3. **(12 points)** You are planning an around-the-world trip with your two best friends for the summer. There is a total of n cities that the three of you want to visit. As you are traveling around the world, you are worried about time zones and airport access. Therefore, some cities can only be visited after visiting another city first, which is in a nearby timezone or has an airport, which are expressed as a list of pairs (cityX,cityY) (cityX can only be visited after visiting cityY).

Given the total number of cities and a list of dependency pairs, is it possible for you all to visit all cities?

Your task is to write the function `can_visit_all_cities`, which determines whether visiting the n cities is possible or not given the dependencies. **(8 points)**

Requirements

- Your code must be written in Python 3, and the filename must be `ps4.py`.
- We will grade only the `can_visit_all_cities` function; please do not change its signature in the starter code. include as many helper functions as you wish.
- You are **not** allowed to use the built-in Python dictionary or set.
- To get full marks, your algorithm must have average-case runtime $\mathcal{O}(m+n)$, where m is the number of dependencies given.

Input/Output Specification Please see the python starter code for input / output specifications.

Write-up (4 points): in your `ps4.pdf/ps4.tex` files, include the following: an explanation of how your code works, justification of correctness, and justification of desired $\mathcal{O}(m+n)$ average-case runtime.